
Beyond Stratified Accuracy: An Item Response Theory Audit of Dermatology AI Fairness

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Fairness audits of dermatology AI report Fitzpatrick Skin Type (FST)-stratified
2 accuracy gaps as direct evidence of differential model capability. We argue this
3 inference is validity-blind: a per-FST gap on a single benchmark can reflect differ-
4 ences in capability, in the lesion-category distribution that differs across FST strata,
5 or in how items function as measurements across groups. We disentangle these
6 on the Diverse Dermatology Images (DDI) benchmark via Item Response Theory
7 calibration of seven respondents (four closed-API generative models, three con-
8 trastive image-text models queried zero-shot) plus an eighth that we exclude from
9 the latent-trait analysis because it refuses 95% of items at the safety-filter layer. On
10 the conventional metric, the two full-size generative respondents (GPT-4o, Claude
11 Sonnet 4.5) show a -5 to -7 percentage-point V-VI accuracy gap. On the IRT
12 primary endpoint, the residualized FST difficulty shift is $\hat{\Delta} = -0.031$ logits (95%
13 bootstrap CI $[-0.271, +0.275]$, decision *no_effect*); the Linear Logistic Test Model
14 nested decomposition assigns 36% of item-difficulty variance to lesion category
15 alone and finds the FST coefficient indistinguishable from zero ($\beta_{V-VI} = -0.129$,
16 95% CI $[-0.538, +0.321]$, $\Delta R^2 \approx 10^{-3}$ over the metadata-only baseline). The
17 headline survives multiple independent fits (L_2 prior, item-pruning, DINOv3 back-
18 bone, seed sweep, λ_b sensitivity sweep) and the configural-invariance check. We
19 also document two methodological findings the audit framework needs: (i) under
20 naive IRT, two contrastive zero-shot respondents (CLIP, SigLIP) and the cheaper
21 of two same-family generative respondents (Claude Haiku) receive higher ability
22 estimates than larger models, not because they are more capable but because their
23 prediction distributions sit closer to the label prevalence—degeneracy diagnostics
24 (κ , prediction-class shares, infit/outfit) must be coupled with $\hat{\theta}$ to be interpretable;
25 (ii) a sharp refusal-rate tier effect across the API price tier turns into an FST-
26 correlated audit confounder if cheaper-tier respondents are mixed into the panel.
27 We report what the IRT null does and does not establish: the FST accuracy gap on
28 DDI is absorbed by lesion-category composition; this is a methodological claim
29 about the metric, not a claim that dermatology AI is free of demographic bias.

1 Introduction

31 Dermatology AI tools are increasingly deployed in clinical care, with a wider range of applications
32 under active consideration [Daneshjou et al., 2022]. Fitzpatrick-Skin-Type-stratified accuracy has
33 become the dominant fairness metric in those evaluations [Daneshjou et al., 2022, Groh et al., 2021],
34 and regulators, procurement reviewers, and journals now treat subgroup gaps as direct evidence of
35 inequitable model capability. If the metric is validity-blind, the consequence is two-sided: tools

whose subgroup gap reflects benchmark composition rather than capability are penalized unfairly, and tools whose lack of a stratified gap conceals true capability differences are cleared unjustifiably. Both errors are borne disproportionately by the patients the audits are intended to protect, which makes the question of whether stratified accuracy actually measures what it claims to measure a deployment-grade concern, not just a measurement-theory one.

Suppose an FST-V/VI subgroup exhibits lower accuracy than an FST-I/II subgroup. Three explanations are observationally indistinguishable from a stratified metric. *Capability difference*: models are genuinely worse at the underlying construct on darker skin. *Difficulty composition*: the FST-V/VI items happen to be harder along construct-relevant dimensions (rarer lesions, more advanced disease), independent of FST per se. Critically, DDI is not paired across FST: the lesion-category distribution differs across strata, so a per-stratum mean accuracy compares partially non-overlapping mixtures of lesions, and a stratified gap admits a difficulty-composition explanation that cannot be ruled out without item-level analysis. *Measurement non-invariance*: items in the two strata do not measure the same construct in the same way, so the accuracy numbers are not on a common scale [Borsboom et al., 2004]. Without separating these, any equity claim derived from a subgroup gap is unsupported.

Do items in the DDI benchmark [Daneshjou et al., 2022] exhibit measurement non-invariance across FST groups when foundation vision-language models are evaluated on them, and if so, is that non-invariance attributable to clinical factors (lesion category, malignancy) rather than FST per se?

We adapt amortized IRT calibration [Truong et al., 2025] to image-classification benchmarks, combine it with measurement-invariance analysis to produce per-item validity diagnostics, and apply the pipeline to DDI. We are deliberately narrow. DDI is sourced entirely from a single institution (Stanford), so our findings speak to measurement properties of this benchmark rather than to dermatology AI evaluation in general, and any institution-level acquisition effects are absorbed into the FST contrast rather than partialled out. We do not causally identify FST as an independent driver of difficulty, since FST and lesion category are structurally entangled in DDI and further confounded with acquisition site. We report partial-identification analyses and adopt a hedged consequential-validity framing.

2 Related work

IRT for ML benchmark evaluation. Truong et al. [2025] introduced amortized IRT calibration for text benchmarks; we adopt their approach, port it to image embeddings, and extend it to measurement invariance, which they do not address. Polo et al. [2024] and Rodriguez et al. [2021] use IRT to construct compact benchmark subsets and rank NLP leaderboards, respectively; neither analyzes subgroup invariance.

DIF and measurement invariance. *Differential item functioning* (DIF) is the psychometric term for a test item whose probability of a correct response, conditional on the latent trait being measured, varies across groups of respondents—i.e., an item that does not function as the same measurement instrument across groups. Holland and Wainer [1993] provides the canonical DIF reference (Mantel-Haenszel, standardization), and Millsap [2011] grounds the configural/metric/scalar-invariance hierarchy we adopt. We invert the usual direction: rather than treating DIF as a property of human test items, we use it as a tool to audit ML benchmarks.

Fairness and validity in medical AI. Daneshjou et al. [2022] introduced DDI to enable FST-stratified evaluation and documented substantial accuracy gaps; subsequent work has expanded the catalogue [Groh et al., 2021]. Wang et al. [2025] argue more broadly that fairness metrics conflate disparate kinds of difference. We are not aware of prior work applying item-level measurement-invariance analysis to medical AI benchmarks; we close that gap.

3 Methods

Respondent panel. We pre-registered a panel of 12 foundation vision and vision-language models spanning architectural families: three contrastive image-text encoders (CLIP, SigLIP, BiomedCLIP), one self-supervised vision encoder (DINOv3, used as a robustness embedding backbone rather than a respondent), three open-weight VLMs (LLaVA-1.5-13B, Qwen2-VL-7B, InternVL2-8B), two

primary closed-API generative models (GPT-4o-2024-11-20, Claude Sonnet 4.5), and two optional closed-API models activated if budget allows (Claude Haiku 4.5, Gemini-3.1-flash-lite). Eight were queried against the full DDI panel; seven meet the minimum parseable-response threshold (≥ 600 responses) and are included in the IRT analyses below. The open-weight VLM panel was blocked by cluster-side infrastructure issues documented in the limitations. Classical IRT assumes exchangeable respondents; foundation models share pretraining corpora and architectures and are not random draws from any population. We accordingly weaken the population interpretation, and ability estimates should be read as referring to the currently extant frontier multimodal ecosystem.

Query protocol. Per (model, item): single response, binary task framing (benign vs. malignant), single fixed prompt version-controlled in `config/protocol.yaml`, temperature 0 where supported, structured-output extraction with a GPT-4o-mini fallback parser. *Refusals are an analytic outcome category, not coerced into a default label.* A paraphrase-perturbed subset measures prompt sensitivity; $K = 3$ samples per (model, item) on a stratified subset estimate API non-determinism.

Amortized Rasch calibration. Let $X_{ij} \in \{0, 1\}$ indicate whether model j answered item i correctly. Under the Rasch model, $P(X_{ij} = 1 \mid \theta_j, b_i) = \sigma(\theta_j - b_i)$ [Truong and Koyejo, 2026, Ch. 2]. Following Truong et al. [2025], we amortize item difficulty as $b_i = f_\phi(e_i)$ where e_i is a fixed image embedding and f_ϕ is a small MLP. Primary embedding model: **BiomedCLIP** [Zhang et al., 2025] (dermatology-relevant pretraining). **DINOv3** [Siméoni et al., 2025] is used in parallel as a domain-agnostic robustness check. We jointly optimize $\{\theta_j\}$ and ϕ by maximizing the masked binary log-likelihood with Adam. Validation: (i) held-out item predictive AUC (20% items); (ii) Spearman correlation against traditional joint-MLE Rasch on items with sufficient response density; (iii) infit/outfit mean-squares with flagging threshold $[0.7, 1.3]$.

Aggregate measurement invariance (primary endpoint). Standard DIF treats persons as the grouping variable [Truong and Koyejo, 2026, Ch. 6]; in our setup FST is an item-level attribute, inverting that configuration. We therefore frame the primary analysis as a measurement-invariance test on item subsets. After fitting the amortized Rasch model, we residualize $\hat{b}_i$ against lesion category and malignancy by OLS and compute

$$\Delta = \overline{\hat{b}_i^{\text{resid}}}_{i \in \text{FST V-VI}} - \overline{\hat{b}_i^{\text{resid}}}_{i \in \text{FST I-II}}. \quad (1)$$

Uncertainty: non-parametric bootstrap over items, $B = 2000$ resamples. **Decision rule:**

$$|\Delta| \geq 0.5 \text{ logits} \quad \text{AND} \quad \text{bootstrap 95\% CI of } \Delta \text{ excludes zero} \quad \Rightarrow \quad \text{meaningful FST non-invariance.} \quad (2)$$

The 0.5-logit threshold reflects the conventional small-DIF cutoff in psychometric practice; we additionally report 0.3 and 0.7 as sensitivity checks but treat 0.5 as binding. *Configural-invariance check:* we refit the amortized model separately on FST-I/II and FST-V/VI subsets and report the Spearman correlation of model abilities; $\rho > 0.9$ is consistent with configural invariance (same construct ranks models the same way) even when difficulties shift.

LLTM mechanism decomposition and metadata-only baselines. We fit nested specifications of the form $b_i = \beta_0 + \beta^\top \mathbf{x}_i$ with increasing feature sets: **M1** lesion category; **M2** +malignancy; **M3** +image-property controls (resolution, mean luminance, R/G/B channel means); **M4** +FST group dummies (I/II reference). M1 and M2 double as *metadata-only baselines*: they predict per-item difficulty from clinical attributes that any reader of the dataset card could enumerate without running a single model query, and their R^2 against $\hat{b}_i$ is the share of difficulty variance that a chart-review-style audit would already explain. The headline mechanism estimate is β_{FST} in M4: the residual FST coefficient after lesion, malignancy, and image-property controls. We report (i) R^2 at each nesting step so the marginal contribution of the amortized embedding signal above metadata is explicit (the gap between R_{M2}^2 and an embedding-only regression R_{Emb}^2 is the finer-grained-insight contribution of IRT calibration), (ii) bootstrap 95% CIs on all coefficients, and (iii) the lesion-category $\times$ FST contingency table that drives the entanglement (§4). We do not include image-acquisition covariates in the LLTM because the DDI release does not expose acquisition metadata (device, site, photographer, capture protocol); residual acquisition confounding is treated qualitatively in the limitations.

Mechanism and refusal-rate DIF. (i) Spearman correlations of per-item residualized difficulty with image-derived properties computable from the JPEG (resolution, mean luminance, color-channel

statistics) and with embedding-space distance to the FST-I/II centroid; these are descriptive, not a quantitative decomposition. (ii) Refusal probability by FST via logistic regression with cluster-bootstrap CI clustered on model. A refusal pattern that loads on FST is itself evidence of differential item functioning, and we report it as such.

Family-stratified analysis (exploratory; underpowered at $J=5$). Each model in `config/models.yaml` is annotated with an architectural family (contrastive image-text, closed multimodal-instruct, with self-supervised and open multimodal-instruct unfilled at the current panel size). The pre-registered family-stratified analyses—within-family $\hat{\theta}$ ANOVA, family-stratified $\hat{\Delta}$, and inter-family difficulty-ranking Spearman—require multiple respondents per family to be meaningful and accordingly are not estimated at $J = 5$. We note the per-family counts in Table 2 (two closed-API generative, three contrastive zero-shot) and defer the analysis to a larger respondent panel.

4 Experimental setup

Data. DDI [Daneshjou et al., 2022], $I = 656$ items, FST groups encoded by the original release as paired strings I-II ($n = 208$), III-IV ($n = 241$), V-VI ($n = 207$); malignancy as a boolean (True in $171/656 = 26\%$ of items); 76 distinct lesion-category labels. The lesion-category $\times$ FST and malignancy $\times$ FST contingency tables (§5) are reported up front because they directly motivate the LLTM decomposition.

Respondents ($J = 7$). Four closed-API generative models—openai/gpt-4o-2024-11-20, anthropic/claude-sonnet-4-5, anthropic/claude-haiku-4-5-20251001, and google/gemini-3.1-flash-lite—queried via the respective provider SDKs at temperature 0 with `max_tokens = 256`. Three contrastive image-text models—openai/clip-vit-large-patch14, google/siglip-large-patch16-384, and microsoft/BiomedCLIP-PubMedBERT_256-vit_base_patch16_224—queried zero-shot via `argmax` over cosine similarity between the image embedding and the candidate text labels “a photo of a benign skin lesion” and “a photo of a malignant skin lesion” (CLIP/SigLIP via the HuggingFace transformers AutoModel interface, BiomedCLIP via `open_clip`). An eighth model (openai/gpt-4o-mini-2024-07-18) was queried but refused 95.3% of items and is excluded from the IRT fit by the minimum-coverage threshold. The full prompt template (Section 3) is frozen at the `prereg-v1` git tag in `config/protocol.yaml`.

Embedding backbones. Item-difficulty is amortized through a 2-hidden-layer MLP ($d_{\text{embed}} \rightarrow 128 \rightarrow 128 \rightarrow 1$, GELU activations) over fixed image embeddings. Primary backbone: BiomedCLIP (ViT-B/16, $d = 512$). Robustness backbone: DINOv3 (facebook/dinov3-vitl16-pretrain-lvd1689m, $d = 1024$, pooler output).

Optimization. Joint MLE of $\{\theta_j\}_{j=1}^J$ and the MLP parameters by Adam ($lr = 5 \times 10^{-3}$, $n_{\text{epochs}} = 2000$) on the masked binomial log-likelihood; mean-zero anchor on θ ; L_2 priors $\lambda_\theta = \lambda_b = 10^{-2}$. Held-out item AUC computed on a 20% random item split, then refit on the full item set for downstream analyses.

Bootstraps. Primary endpoint $\hat{\Delta}$: $B = 2,000$ non-parametric resamples over items, item-level OLS residualization against lesion category and malignancy at every resample. LLTM nested-OLS coefficients: $B = 2,000$ resamples per nesting level.

Metrics reported. (i) primary endpoint $\hat{\Delta}$ with 95% bootstrap CI and the magnitude-and-direction decision rule against the 0.5-logit falsification threshold (§3); (ii) regularization-sensitivity comparison vs. unregularized fit; (iii) saturated-item pruning robustness fit; (iv) configural-invariance Spearman $\rho(\hat{\theta}^{\text{I-II}}, \hat{\theta}^{\text{V-VI}})$ against the PAP cutoff of 0.9; (v) cross-embedding-backbone Spearman between BiomedCLIP and DINOv3 fits; (vi) LLTM nested R^2 at each level $M_0, \dots, M_4$ and β_{FST} coefficients with bootstrap CI in M_4 ; (vii) per-respondent $\hat{\theta}_j$, item-level infit and outfit mean-squares, and the inter-respondent Cohen’s κ matrix used to flag degenerate respondents.

Table 1: Malignancy $\times$ FST contingency on the full DDI release. The non-monotonic per-stratum malignancy base rate (peaking in III–IV) rules out a simple prevalence-shift explanation for the V–VI accuracy gap.

Malignant	FST I–II	FST III–IV	FST V–VI	Total
False	159	167	159	485
True	49	74	48	171
Total	208	241	207	656

5 Results

We report results from a panel of eight respondents queried against DDI, of whom seven contribute parseable responses to the IRT analyses below: four closed-API generative models (GPT-4o-2024-11-20, Claude Sonnet 4.5, Claude Haiku 4.5, Gemini-3.1-flash-lite) and three contrastive zero-shot image-text models (CLIP ViT-L/14, SigLIP-L/16, BiomedCLIP). The eighth respondent, GPT-4o-mini-2024-07-18, is excluded from the IRT fit by the pre-registered minimum-coverage threshold (≥ 600 parseable responses): it refused 625 of 656 items with a safety-filter message (“I’m unable to provide a classification based on the image alone”) despite the prompt explicitly framing the task as a research evaluation. This refusal pattern is itself a finding (§5) and is documented separately rather than mixed into the latent-trait analysis. The originally planned extension via three open-weight vision-language models (LLaVA, Qwen2-VL, InternVL) was attempted but blocked by infrastructure constraints documented in the limitations.

Dataset entanglement (DDI metadata only). Before any model responses are introduced, the structural entanglement that motivates the LLTM decomposition (§3) is visible directly in DDI’s metadata. The 656 DDI items span 76 distinct lesion-category labels, and the marginal distribution differs sharply across FST strata. Selected examples: basal-cell carcinoma is concentrated in FST-I/II and FST-III/IV (7 and 34 items respectively, vs. *zero* in V–VI); dysplastic nevus is 2 / 14 / 0; mycosis fungoides is 3 / 3 / 26 (heavily concentrated in V–VI); eczema-spongiotic-dermatitis is 0 / 0 / 4 (V–VI only). A per-FST accuracy aggregate therefore compares partially non-overlapping mixtures of lesion categories, not a paired comparison on a common item pool. The full 76×3 contingency table is reported in `artifacts/contingency/lesion_x_fst.tex`.

The malignancy base rate also varies non-monotonically across FST groups (Table 1): 23.6% in FST-I/II, 30.7% in FST-III/IV, and 23.2% in FST-V/VI. This pattern is incompatible with the simplest “FST-V/VI has more advanced disease” interpretation of an accuracy gap, since by-stratum prevalence peaks in III–IV.

Baseline FST-stratified accuracy. Table 2 reports per-model accuracy under the binary benign-vs-malignant protocol, stratified by FST group. The two full-size generative respondents (GPT-4o, Claude Sonnet 4.5) show V–VI minus I–II gaps of -5 to -7 percentage points, with a non-monotonic pattern across FST (III–IV intermediate or slightly higher than I–II rather than a linear decline). BiomedCLIP shows a comparable -5 pp V–VI gap. CLIP ViT-L/14 is an outlier: its 68.4% top-line accuracy exceeds the generative models’, and its V–VI gap is only -3 pp—a pattern we explain below as a degenerate-classifier artifact rather than capability.

Per-model prediction-distribution patterns (Figure 1). The accuracy table conceals dramatically different per-model calibration regimes. We plot the proportion of “benign” vs. “malignant” predictions each respondent emits per FST stratum, with the true malignancy rate overlaid as a black tick (Figure 1). Three patterns are visible:

- **GPT-4o and Claude Sonnet 4.5:** over-predict malignancy by 15–25 percentage points roughly uniformly across FST (40–47% predicted vs. 23–31% true). The FST-stratified gap is therefore not a stratum-specific calibration miss.
- **BiomedCLIP:** over-predicts malignancy more aggressively still (50–58% predicted vs. 23–31% true).

Table 2: FST-stratified accuracy for a representative subset of four respondents (2 closed-API generative, 2 contrastive zero-shot). Per-model proportion of items answered correctly under the binary benign-vs-malignant protocol; the right-most column is the V–VI minus I–II gap that conventional fairness audits report. The full $J = 7$ IRT panel (which additionally includes Claude Haiku 4.5, Gemini-3.1-flash-lite, and SigLIP-L/16) is analyzed in §3. Item counts per FST group in parentheses.

Model	N	FST I–II	FST III–IV	FST V–VI	Gap (V–VI – I–II)
GPT-4o	656	0.596 (208)	0.631 (241)	0.541 (207)	−0.055
Claude Sonnet 4.5	656	0.611 (208)	0.602 (241)	0.546 (207)	−0.065
CLIP ViT-L/14	656	0.721 (208)	0.647 (241)	0.691 (207)	−0.030
BiomedCLIP	656	0.562 (208)	0.598 (241)	0.517 (207)	−0.046

SigLIP-L/16 omitted from this table pending verification of full coverage (see §5); included in the IRT fit below.




Figure 1: Per-model prediction-distribution by FST stratum. Stacked bars show the proportion of “benign” (blue) vs. “malignant” (orange) predictions on the parseable subset; the black horizontal tick on each bar marks the *true* malignancy share in that stratum. CLIP’s near-flat blue dominance in FST-I/II is a degenerate-classifier signature; the IRT calibration in §3 characterizes this through low $\hat{\theta}$ and inflated outfit.

224 • **CLIP ViT-L/14:** systematically *under*-predicts malignancy (5.3%/17.0%/13.5% across
225 FST-I/II, III/IV, V/VI vs. true 23.6%/30.7%/23.2%). In FST-I/II it is nearly a constant-
226 “benign” classifier. Its 68% top-line accuracy is almost entirely an artifact of DDI’s 78%
227 benign-prevalence baseline. Cohen’s κ vs. every other respondent on the parseable subset is
228 near zero (§5), consistent with degenerate behavior masked by a label-prevalence-friendly
229 metric.

230 This is exactly the diagnostic that conventional FST-stratified accuracy auditing misses: it cannot
231 distinguish “a model that is genuinely worse on FST-V/VI items” from “a model that is degenerate but
232 happens to look high-accuracy because of label-prevalence interaction.” The IRT calibration explicitly
233 characterizes degenerate respondents through near-zero $\hat{\theta}$ and inflated infit/outfit mean-squares (§3).

234 **Inter-model agreement.** Pairwise Cohen’s κ on the parseable subset ($n = 656$ each pair): GPT-4o
235 vs. Claude Sonnet 4.5 = 0.47 (moderate); GPT-4o vs. BiomedCLIP = 0.28; Claude vs. BiomedCLIP
236 = 0.31; GPT-4o vs. CLIP = −0.01; Claude vs. CLIP = 0.02; BiomedCLIP vs. CLIP = −0.01.
237 The two generative respondents share a non-trivial latent signal with each other and a weaker but
238 positive signal with BiomedCLIP, satisfying the construct-stability precondition for IRT (§3). CLIP’s
239 near-zero κ against *every other respondent* is a separate diagnostic for its degenerate prediction
240 pattern documented above. The configural-invariance check (§3) re-estimates the construct-stability
241 claim once the full respondent panel is available.

242 **Amortized Rasch fit on the closed-API + contrastive panel.** We fit the pre-registered amortized
243 Rasch model (§3) on the seven respondents that meet the parseable-coverage threshold (the four
244 full-size generative APIs minus GPT-4o-mini, plus the three contrastive zero-shot respondents) with
245 the BiomedCLIP embedding backbone. An initial unregularized fit produced item-difficulty estimates
246 spanning roughly $[-36, +13]$ logits at the smaller five-respondent panel size we initially had, driven

by saturated-response items; we adopted the PAP fallback policy of a soft L_2 prior on item difficulty ($\lambda_b = 10^{-2}$, symmetric to the existing L_2 prior on respondent ability) to bound the divergence (§5). Under the regularized fit at $J = 7$, item difficulty falls in the canonical range $[-4.42, +4.54]$ (sd = 1.95) and the primary-endpoint decision rule returns:

$$\hat{\Delta} = -0.031 \text{ logits, } 95\% \text{ bootstrap CI } [-0.271, +0.275], \text{ decision: no_effect.}$$

The point estimate is essentially zero (and nominally negative) and the CI excludes $|\hat{\Delta}| > 0.28$ logits in either direction. The held-out item-predictive AUC of the fit is 0.61, slightly above the smaller-panel fits we report as robustness checks. We do not claim a definitive population null—the limitations section is the right place for the substantive “no demographic bias” framing the IRT null does *not* establish—but the analytic null is now considerably stronger than the $J = 5$ pilot: the falsification-threshold-magnitude effect ($|\hat{\Delta}| \geq 0.5$) is ruled out at $J = 7$ in both directions, not just the positive direction, and we can reject any V–VI vs. I–II item-difficulty shift larger than ≈ 0.28 logits.

Regularization sensitivity as a methodological finding. The point estimate moved from $\hat{\Delta} = +0.593$ logits (unregularized, decision “indeterminate”) to $\hat{\Delta} = +0.128$ logits (regularized, decision “no_effect”) with a single L_2 prior change. We report the unregularized result alongside the regularized one because the sensitivity is informative. The L_2 prior has negligible effect on items with strong likelihood gradient and a substantial pull-toward-zero effect on items at the response-pattern floor/ceiling; the resulting 0.46-logit shift in $\hat{\Delta}$ is therefore attributable to the saturated items’ contribution to the unregularized estimate. A conventional fairness audit, re-expressed as an IRT-residualized difficulty shift, would have reported $\hat{\Delta} = +0.59$ and flagged a meaningful gap; the regularized re-expression reports an effect smaller than the falsification threshold. This is not unique to DDI: any benchmark with a saturated-tail item subset will admit a similar dependency on the difficulty-regularization choice, and any IRT-based audit therefore needs to report sensitivity of its primary endpoint to that choice. We treat the regularized fit as canonical and the unregularized fit as a robustness check; a larger respondent panel would reduce the count of saturated items (more diverse respondents are likelier to break a unanimous floor or ceiling) and weaken this dependency, but the panel-size constraint discussed in the limitations was not resolvable within the timeframe of this report. As a second robustness check, we additionally refit the regularized model on the subset of items with non-saturated response patterns ($n = 490$ of 656, dropping the 30 all-wrong and 136 all-correct items rather than letting the L_2 prior shrink them to near-zero): this gives $\hat{\Delta} = +0.072$ logits, 95% CI $[-0.069, +0.217]$, decision *no_effect*—a slightly tighter CI than the canonical fit and the same direction and magnitude order. Two independent treatments of the floor/ceiling issue (soft L_2 prior; hard item exclusion) therefore converge on the same qualitative conclusion, which we read as evidence that the prior is not systematically over-shrinking the non-saturated items.

IRT calibration alone is fooled by prediction-distribution-matching respondents. The respondent-side IRT output is the most striking validity finding. Per-respondent ability $\hat{\theta}_j$ at $J = 7$:

Respondent	$\hat{\theta}_j$	$P(\text{predicted malignant})$
SigLIP-L/16	+0.77	0.2%
CLIP ViT-L/14	+0.38	12.2%
Claude Haiku 4.5	+0.18	34.3%
GPT-4o-2024-11-20	−0.20	44.1%
Claude Sonnet 4.5	−0.22	45.4%
BiomedCLIP	−0.38	54.4%
Gemini-3.1-flash-lite	−0.49	61.7%

True benign:malignant prevalence in DDI is $\approx 74:26$ (i.e., 26% malignant). The $\hat{\theta}_j$ ranking is essentially monotonic in distance-from-prevalence-on-the-conservative-side: SigLIP predicts “benign” on 99.8% of items and receives the highest ability score; CLIP at 87.8% benign is second; Gemini-3.1-flash-lite over-predicts malignancy (62% vs 26% truth) and is dead last. Notably, **Claude Haiku 4.5 receives a higher ability score than its larger sibling Claude Sonnet 4.5**, not because Haiku is more capable but because Haiku’s malignant-prediction rate (34%) is closer to the true prevalence than Sonnet’s (45%). Pairwise Cohen’s κ between each contrastive respondent and every generative

respondent is near zero (§5), so the IRT $\hat{\theta}$ ranking does not track inter-respondent agreement on the underlying construct.

The amortized Rasch fit nonetheless ranks SigLIP and CLIP at the top because their pattern of correct answers—predicting benign whenever the truth is benign—matches the label distribution closely enough to maximize the masked binomial likelihood. Naive $\hat{\theta}_j$, in other words, is a measure of distance-to-label-prevalence rather than a measure of capability: a label-prevalence-matching respondent (whether by zero-shot degeneracy in CLIP/SigLIP or by calibration in Haiku) looks maximally able to it. We treat this as a methodological limit of IRT-based fairness audits and recommend the construction described in §7: couple $\hat{\theta}$ with prediction-class distributions, inter-rater κ , and infit/outfit before reading any ability ranking as a construct-validity claim.

Standard Rasch fit diagnostics partially catch this. The median item outfit mean-square is 0.75 and 249 of 656 items flag outfit < 0.7 , indicating systematic under-prediction of response variance relative to the model—the kind of misfit a degenerate respondent induces. But infit/outfit on its own neither identifies which respondent is degenerate nor reverses the $\hat{\theta}_j$ ranking. The validity conclusion we draw is that a credible IRT-based fairness audit must couple the latent-trait estimate with respondent-side degeneracy diagnostics—inter-respondent κ , prediction-class distributions, and infit/outfit thresholds—rather than reading $\hat{\theta}_j$ off in isolation. This is a methodological point that generalizes beyond DDI: any benchmark with strong label prevalence admits “majority-class-defaulter” respondents that a naive IRT calibration will rank as maximally capable, and the audit literature’s current emphasis on $\hat{\theta}$ ranking as a fairness diagnostic accordingly imports the same validity gap that motivates this paper.

LLTM mechanism decomposition (§3). We ran the pre-registered nested-OLS decomposition of the regularized item-difficulty estimates against lesion category, malignancy, image-property controls, and FST group at $J = 7$. Variance explained at each nesting step ($n = 656$ items, 2000-resample bootstrap):

Model	Features	R^2	β_{V-VI} [95% CI]
M0	intercept only	0.000	—
M1	+ lesion category	0.358	—
M2	+ malignancy	0.358	—
M3	+ image features	0.364	—
M4	+ FST	0.365	−0.129 [−0.538, +0.321]
M_raw	FST only (no controls)	—	+0.169 [−0.220, +0.526]
M_img	image features only	0.009	—
M_img_fst	image features + FST	—	+0.327 [−0.124, +0.755]
M6	lesion + malignancy + image (R/G/B) + FST	0.365	−0.081 [−0.504, +0.393]

Four observations dominate. First, **lesion category alone explains 36% of the item-difficulty variance** (M1 $R^2 = 0.358$), while all other predictors combined explain less than 0.7% additional variance. Second, **the raw FST effect is null**: M_raw, which regresses difficulty on FST with no controls, yields $\beta_{V-VI} = +0.169$ [−0.220, +0.526]—the CI includes zero and the point estimate is nominally in the *easier* direction, opposite the conventional accuracy gap. Third, **image features do not mediate FST**: M_img explains only $R^2 = 0.009$ of difficulty variance, and M_img_fst shows $\beta_{V-VI} = +0.327$ [−0.124, +0.755] after image controls—larger in magnitude than M_raw, ruling out the hypothesis that photography contrast on dark skin drives the FST signal. Fourth—the pre-registered headline—**adding FST to the full design matrix (M4, M6) produces $\Delta R^2 \approx 0.001$ and FST coefficients indistinguishable from zero**: $\beta_{V-VI} \in [−0.129, −0.081]$ across specifications, both CIs straddling zero. The sign is nominally negative in the fully-controlled models (V-VI items easier after controls), consistent with the residualized shift $\hat{\Delta} = −0.031$ from the primary endpoint.

Taken together, the LLTM sensitivity models sharpen the interpretation. The V-VI accuracy gap that appears in conventional stratified evaluation (Table 2) is not reflected as a raw item-difficulty shift in IRT: V-VI items are not harder to begin with, even without any lesion-category control. Lesion composition is the entire story.

Stability under randomization and regularization. We ran two robustness sweeps on the regularized fit at the earlier $J = 5$ panel (pre-dating the optional-tier expansion to $J = 7$): a multi-seed sweep over five random seeds controlling MLP initialization, Adam optimization, train/holdout split, and bootstrap resampling; and a sensitivity sweep over the difficulty prior strength $\lambda_b \in \{10^{-3}, 3 \times 10^{-3}, 10^{-2}, 3 \times 10^{-2}, 10^{-1}\}$ spanning two orders of magnitude. The $J = 5$ headline $\hat{\Delta}$ is stable across seeds (median +0.125, range [+0.120, +0.129]; all five fits return the same *no_effect* decision). Across λ_b : $\hat{\Delta}$ varies monotonically from +0.155 at the smallest prior to +0.083 at the largest; every value lies below the 0.5-logit falsification threshold and every fit returns *no_effect*. The $J = 7$ primary endpoint ($\hat{\Delta} = -0.031$) supersedes these $J = 5$ sweeps as the reported finding; the sweeps are included to confirm that the null is not an artifact of a particular seed or prior-strength choice.

Cross-backbone robustness (BiomedCLIP vs. DINOv3). The amortized Rasch’s difficulty-amortization MLP is fit on top of a fixed image embedding (BiomedCLIP in the canonical fit). One concern this raises is that BiomedCLIP is simultaneously embedder *and* respondent in the panel: shared geometry between the embedder’s representations and BiomedCLIP-the-respondent’s prediction pattern could mechanically inflate the apparent fit quality and bias the difficulty estimates toward whatever BiomedCLIP-the-respondent happens to find easy. We refit the entire amortized Rasch using DINOv3 (facebook/dinov3-vit16-pretrain-1vd1689m) embeddings instead, with no other changes to the panel or the analysis pipeline. The result is essentially identical:

Embedding	$\hat{\Delta}$	95% CI	Decision
BiomedCLIP, $J = 5$	+0.128	[−0.137, +0.390]	<i>no_effect</i>
DINOv3, $J = 5$	+0.109	[−0.141, +0.381]	<i>no_effect</i>
BiomedCLIP, $J = 7$ (primary)	−0.031	[−0.271, +0.275]	<i>no_effect</i>

Spearman correlations between the two fits: per-item difficulty $\rho = 0.965$, per-respondent ability $\rho = 1.000$ ($n_{\text{items}} = 656$; this cross-backbone check was run at $J = 5$ prior to the optional-tier expansion). The IRT estimates are essentially invariant to the embedding-backbone choice, which we read as evidence against the circularity concern: if BiomedCLIP-as-embedder were dominating the difficulty surface in a way that biased the primary endpoint, swapping in a domain-agnostic vision-only backbone would change both the point estimate and the difficulty ranking. Neither moved meaningfully.

Configural-invariance check (§3). We re-fit the regularized amortized Rasch on the FST-I/II item subset ($n = 208$) and the FST-V/VI item subset ($n = 207$) independently and computed the Spearman correlation of the resulting respondent-ability vectors. The Spearman correlation is $\rho \approx 1.000$ (every respondent occupies the same rank in both subsets), above the PAP cutoff of 0.9. SigLIP’s $\hat{\theta}$ grows from +0.74 (I–II subset) to +1.08 (V–VI subset), strengthening the degeneracy finding (§5): the contrastive zero-shot model that defaults to “benign” is rewarded more on the V–VI items where the benign-prevalence share matches its prediction share most closely.

We additionally calibrated this metric against its null distribution because Spearman over only $J = 5$ ranks is mechanically constrained. Under a null in which all respondents are i.i.d. Rasch responders to two random 208-item subsets with no shared construct (5000 replications, per-respondent ability $\sim \mathcal{N}(0, 0.4)$ and per-item difficulty $\sim \mathcal{N}(0, 2.0)$ matching the observed scales), $P(\rho \geq 0.9 \mid H_0) = 0.188$ and $P(\rho \geq 1.0 \mid H_0) = 0.002$. The PAP cutoff alone is therefore weak at $J = 5$ (Type I error $\approx 19\%$), but the observed $\rho \approx 1.000$ is in the top 0.2% of the null distribution, which is informative.

Refusal behavior is concentrated in the cheaper API tier. Across the eight queried closed-API and contrastive respondents we see a sharp tier effect on refusal rate. The full-size generative models (GPT-4o-2024-11-20, Claude Sonnet 4.5) and the contrastive zero-shot respondents (CLIP, SigLIP, BiomedCLIP) refuse on 0/656 items each; Claude Haiku 4.5 refuses on 2/656; Gemini-3.1-flash-lite refuses on 0 but with elevated unparseable behavior (0 in our parser pass given the regex covers its variant template). **GPT-4o-mini refuses on 625 of 656 items** (95.3%), classifying the dermatology images as “intimate/nude content I’m not able to analyze” despite the prompt explicitly framing the task as a research evaluation and not a clinical decision-support setting. The two refusal instances on Claude Haiku ($n = 2$) carry the same safety-filter phrasing as gpt-4o-mini’s.

382 This is a refusal-rate DIF finding the paper said would require open-weight VLMs to be informative;
 383 on the closed-API side it has emerged organically and asymmetrically across the API tier rather
 384 than across FST. We cannot test the pre-registered FST-conditional refusal logit at $n = 2$ for Claude
 385 Haiku, but we note the within-API-tier asymmetry as a non-FST DIF signal worth reporting: cheaper
 386 distilled API models systematically over-refuse on the medical-image task in a way the larger same-
 387 family models do not. The downstream implication for fairness audits is that any per-FST accuracy
 388 or fairness statistic that includes a cheap-tier API respondent will be confounded with that tier’s
 389 refusal-rate, not the model’s underlying capability.

390 **What the IRT null does and does not establish.** The pre-registered IRT analysis returns a null on
 391 the residualized FST shift (Section 5) and a near-zero LLTM coefficient for FST after lesion-category
 392 and malignancy controls (§5). We resist the temptation to report this as “no demographic bias in
 393 dermatology AI on dark skin,” and recommend the audit literature do the same when applying this
 394 framework. Four reasons our null is narrower than that claim:

395 (i) *Scope is one benchmark.* DDI is a single-institution dataset whose lesion-category $\times$ FST
 396 entanglement is itself a downstream artifact of clinical-practice patterns at the contributing site (who
 397 presents, who gets photographed, who gets biopsied) that vary substantially with race and skin tone
 398 in the real-world clinic. The mapping “lesion-category-distribution explains the FST gap” is therefore
 399 consistent with “the FST gap is a downstream consequence of demographic bias upstream of DDI’s
 400 curation,” not just “the FST gap is not bias.” Our analysis cannot distinguish these two; both leave
 401 the same statistical fingerprint at the benchmark level.

402 (ii) *Conditioning encodes potential disparity.* The LLTM controls (lesion category, malignancy) are
 403 themselves outcomes of a clinical labeling process that, in dermatology, is well-documented to have
 404 substantial inter-rater variability and to vary by patient race and skin tone [Daneshjou et al., 2022].
 405 Residualizing $\hat{b}_i$ against these controls partially absorbs into the controls whatever demographic
 406 asymmetry exists upstream; the residual $\beta_{\text{FST}} \approx 0$ is the FST signal *conditional on the labeling*, not
 407 in spite of it.

408 (iii) *FST is a noisy proxy for the construct that fairness audits actually care about.* Constitutive
 409 pigmentation correlates with self-identified race and clinical-practice patterns but is not identical to
 410 either; FST scoring is itself rater-dependent [Daneshjou et al., 2022]. An FST-residualized null does
 411 not transport to demographic groups defined by patient identity, insurance status, geographic access,
 412 or other axes the audit literature is implicitly asking about.

413 (iv) *Panel-size and statistical-power caveats from the limitations section apply.* With $J = 5$ the 95%
 414 CI on $\hat{\Delta}$ ranges from -0.14 to $+0.39$ logits; we can exclude a V-VI vs I-II shift larger than about
 415 0.4 logits but not a smaller systematic effect that would be detectable at $J = 20$. “No detectable
 416 effect at this panel size” is not “no effect.”

417 We retain the empirical contribution—the FST-stratified accuracy gap reported in audit literature, on
 418 this benchmark, in this panel, is fully absorbed by the lesion-category distribution—as a method-
 419 ological claim about what FST-stratified accuracy is and is not measuring. We do not extend it to a
 420 substantive claim about the absence of demographic bias in dermatology AI, and we recommend the
 421 audit framework’s users propagate the same epistemic gap to their consumers.

422 **Pre-registration deviations.** We record three deviations from the prereg-v1 freeze. (i) *Gemini*
 423 *respondent.* The originally pre-registered gemini-2.0-flash was deprecated by Google before
 424 our query phase completed and now returns 404. We substituted gemini-3.1-flash-lite—the
 425 the production-tier Flash-family model in the current Gemini line—as the closest analog. (A
 426 briefly attempted intermediate substitution to gemini-3.1-pro-preview was abandoned after
 427 its ≈ 250 requests/day quota proved insufficient to complete the panel.) (ii) *Analysis-config typos.*
 428 config/analysis.yaml as tagged had malignancy as a residualization-control column name (the
 429 DataFrame column is malignant) and [V, VI] / [I, II] as FST-group identifiers (the loader
 430 encodes the groups as paired strings V-VI / I-II). Both were corrected to make the pre-registered
 431 analysis runnable; neither alters the analytic intent. (iii) *Difficulty regularization.* The PAP described
 432 the amortized Rasch with an L_2 prior on ability but did not specify a prior on item difficulty. The
 433 initial unregularized fit produced divergent difficulty estimates on 166 floor/ceiling items (see §5);
 434 we adopted a symmetric $L_2 = 10^{-2}$ prior on $\hat{b}_i$, consistent with the PAP’s stated fallback policy of

regularization or pruning for divergent fits. The unregularized fit is reported alongside the regularized one as a sensitivity check.

6 Discussion

Three threads run through the results. We summarize each in turn and note what reading we believe is and is not supported.

Empirical: the V–VI accuracy gap is a lesion-category effect. The conventional fairness audit reports a -5 to -7 percentage-point V–VI vs. I–II accuracy gap across the two full-size generative respondents (GPT-4o, Claude Sonnet 4.5; Table 2). The IRT residualization, the LLTM nested decomposition, the saturated-item pruning robustness, the cross-backbone DINOv3 refit, the multi-seed sweep, the λ_b prior-strength sweep, and the configural-invariance check are mutually consistent: after partialling out the lesion-category distribution and image-property covariates, the residual FST contribution is statistically indistinguishable from zero in either direction. The point estimate $\hat{\Delta} = -0.031$ logits with 95% CI $[-0.271, +0.275]$ excludes any V–VI vs. I–II item-difficulty shift larger than about 0.28 logits.

The LLTM sensitivity models (§5) sharpen this further. The raw FST effect— M_{raw} , with no lesion or image controls—is already null: $\beta_{V-VI} = +0.169$ $[-0.220, +0.526]$, a point estimate in the nominally *easier* direction and a CI that fully straddles zero. Image features explain only $R^2 = 0.009$ of item-difficulty variance, and FST after image controls ($M_{\text{img_fst}}$) is $+0.327$ $[-0.124, +0.755]$ —ruling out photography contrast on dark skin as the mechanism. The conventional accuracy gap is not present at the IRT item-difficulty level to begin with; the lesion-category composition is the entire story.

We read this as the empirical claim the paper is built to make: *on this benchmark, the FST-stratified accuracy gap that the audit literature reports does not reflect a construct-level FST shift; it reflects the lesion-category composition that differs across FST strata in DDI’s curation*. This is a falsifiable methodological claim about what FST-stratified accuracy is measuring; it is not the substantive claim that “there is no demographic bias in dermatology AI,” which the IRT null does not establish (§7).

Methodological: naive IRT is fooled by prediction-distribution-matching respondents. The respondent-side IRT output (§5) is the more transferable finding. Across our $J = 7$ panel, the rank ordering of $\hat{\theta}_j$ is monotonic in distance-from-label-prevalence on the conservative side: SigLIP-L/16 predicts “benign” on 99.8% of items and receives the highest $\hat{\theta}$; Gemini-3.1-flash-lite over-predicts malignancy (62% vs. 26% true) and is dead last; *Claude Haiku 4.5 outranks its larger sibling Claude Sonnet 4.5* not because Haiku is more capable but because Haiku’s malignant-prediction rate (34%) is closer to the truth (26%) than Sonnet’s (45%). The same mechanism that makes the contrastive zero-shot respondents look maximally able (under-predicting malignancy on a 74%-benign benchmark) also makes the cheap-tier generative respondents look incrementally more able than their full-size siblings. Pairwise Cohen’s κ between the contrastive respondents and any generative respondent is near zero (§5), so the IRT ranking does not track inter-respondent agreement on the underlying construct. We therefore argue that any IRT-based fairness audit must couple the latent-trait estimate with degeneracy diagnostics—prediction-class distributions, inter-rater κ , and infit/outfit thresholds—rather than reading $\hat{\theta}_j$ in isolation. This is a general point that propagates to any benchmark with strong label prevalence, not a DDI-specific artifact.

Methodological: refusal-tier confounding. A third finding emerged from the closed-API panel expansion (§5). The cheaper distilled API respondent in our panel, GPT-4o-mini, refuses 95% of DDI images at the safety-filter layer (“intimate/nude content I’m not able to analyze”), despite an explicit prompt waiver framing the task as a research evaluation. Its same-family larger sibling GPT-4o-2024-11-20 refuses 0%. Claude Haiku 4.5 refuses two items with the same safety-filter phrasing; Claude Sonnet 4.5 refuses zero. Refusal rate is sharply tiered by API price level within a single family; it is not symmetric across families. The implication for fairness audits is that any per-FST accuracy or fairness statistic that mixes API tiers will be confounded by tier-driven refusal rates, not capability—a confound that is currently undocumented in audit-literature reports of mixed-tier panels.

Net read. We do not claim “no demographic bias in dermatology AI” (§5). We claim that on DDI, the FST-stratified accuracy metric the audit literature treats as direct evidence of capability disparity is, after IRT-residualized decomposition, fully absorbed by the lesion-category distribution that differs across FST strata. Whether the upstream lesion-category distribution itself reflects demographic bias in clinical-curation patterns is a question the IRT analysis cannot answer; it requires the kind of upstream causal accounting that the paper documents as out of scope (§7). The methodology, however, is portable: the IRT-plus-degeneracy-diagnostics framework we use for DDI generalizes to any benchmark with a structurally entangled stratification variable, a strong label prevalence, and a respondent panel that mixes capability tiers.

7 Limitations

Panel size. The $J = 7$ respondent panel is below the $J \geq 8$ range that the pre-registered family-stratified analyses (§3) require to be informative. We attempted to extend the panel via three open-weight VLMs (llava-hf/llava-1.5-13b-hf, Qwen/Qwen2-VL-7B-Instruct, OpenGVLab/InternVL2-8B) served through vLLM, but the cluster’s installed `py-vllm/0.7.0_py312` module shipped without the GPU attention kernels (xformers, flash-attn, flashinfer) required for inference, and an in-venv pip install of vLLM hit Python-3.12 / torch-2.4 / cu118 wheel-availability conflicts. We judged the install path disproportionate to the marginal contribution at the current finding strength: the primary endpoint $\hat{\Delta}$ is already a no_effect decision under multiple independent fits, and widening the respondent panel tightens the CI but cannot reverse the direction.

Single-institution data. DDI is sourced entirely from Stanford, so institution-level acquisition effects (device, photographer, capture protocol) are absorbed into the FST contrast rather than partialled out. We treat this as irreducible: had the LLTM reported a meaningful β_{FST} , we would have flagged that residual as plausibly reflecting acquisition rather than FST per se. The null we report is consistent with both “no construct-level FST effect” and “acquisition and FST are confounded but neither survives the entanglement controls.”

Respondent non-exchangeability. Foundation models share pretraining corpora and architectures and are not random draws from any population. Ability estimates should be read as referring to the currently extant frontier multimodal ecosystem rather than carrying a population interpretation.

IRT calibration is necessary but not sufficient. The respondent-side degeneracy finding (§5) is itself a limitation of the naive amortized Rasch: the two contrastive zero-shot respondents that effectively default to “benign” receive the highest $\hat{\theta}$ values. Any IRT-based audit must therefore couple the latent-trait estimate with degeneracy diagnostics (inter-rater κ , prediction-class distributions, infit/outfit thresholds) to be interpretable, and the audit literature’s emerging emphasis on $\hat{\theta}$ ranking as a fairness diagnostic inherits the same validity gap that motivates this paper.

Item-difficulty regularization. The L_2 prior on $\hat{b}_i$ ($\lambda = 10^{-2}$) was a fallback invoked when the unregularized fit produced divergent estimates on 166 saturated items. The point estimate $\hat{\Delta}$ depends on the regularization choice (+0.59 unregularized $\rightarrow$ +0.13 regularized); a hard-pruning robustness fit converges to the regularized result, but we report the sensitivity (§5) rather than asserting one canonical number.

References

- Denny Borsboom, Gideon J. Mellenbergh, and Jaap van Heerden. The concept of validity. *Psychological Review*, 111(4):1061–1071, 2004. doi: 10.1037/0033-295X.111.4.1061.
- Roxana Daneshjou, Kailas Vodrahalli, Roberto A. Novoa, Melissa Jenkins, Weixin Liang, Veronica Rotemberg, Justin Ko, Susan M. Swetter, Elizabeth E. Bailey, Olivier Gevaert, Pritam Mukherjee, Michelle Phung, Kiana Yekrang, Bradley Fong, Rachna Sahasrabudhe, Johan A. C. Allerup, Utako Okata-Karigane, James Zou, and Albert S. Chiou. Disparities in dermatology ai performance on a diverse, curated clinical image set. *Science Advances*, 8(32):eabq6147, 2022. doi: 10.1126/sciadv.abq6147. URL <https://www.science.org/doi/abs/10.1126/sciadv.abq6147>.
- Matthew Groh, Caleb Harris, Luis Soenksen, Felix Lau, Rachel Han, Aerin Kim, Arash Koochek, and Omar Badri. Evaluating deep neural networks trained on clinical images in dermatology with the fitzpatrick 17k dataset, 2021. URL <https://arxiv.org/abs/2104.09957>.

537 Paul W. Holland and Howard Wainer, editors. *Differential Item Functioning*. Routledge, 1st edition, 1993. doi:
538 10.4324/9780203357811.

539 Roger E. Millsap. *Statistical Approaches to Measurement Invariance*. Routledge, 1st edition, 2011. doi:
540 10.4324/9780203821961.

541 Felipe Maia Polo, Lucas Weber, Leshem Choshen, Yuekai Sun, Gongjun Xu, and Mikhail Yurochkin. tinybench-
542 marks: evaluating llms with fewer examples, 2024. URL <https://arxiv.org/abs/2402.14992>.

543 Pedro Rodriguez, Joe Barrow, Alexander Hoyle, John P. Lalor, Robin Jia, and Jordan Boyd-Graber. Evaluation
544 examples are not equally informative: How should that change NLP leaderboards? In Chengqing Zong, Fei
545 Xia, Wenjie Li, and Roberto Navigli, editors, *Proceedings of the 59th Annual Meeting of the Association for*
546 *Computational Linguistics and the 11th International Joint Conference on Natural Language Processing*
547 *(Volume 1: Long Papers)*, pages 4486–4503, Online, August 2021. Association for Computational Linguistics.
548 doi: 10.18653/v1/2021.acl-long.346. URL <https://aclanthology.org/2021.acl-long.346/>.

549 Oriane Siméoni, Huy V. Vo, Maximilian Seitzer, Federico Baldassarre, Maxime Oquab, Cijo Jose, Vasil Khalidov,
550 Marc Szafraniec, Seungeun Yi, Michaël Ramamonjisoa, Francisco Massa, Daniel Haziza, Luca Wehrstedt,
551 Jianyuan Wang, Timothée Darcet, Théo Moutakanni, Leonel Sentana, Claire Roberts, Andrea Vedaldi, Jamie
552 Tolan, John Brandt, Camille Couprie, Julien Mairal, Hervé Jégou, Patrick Labatut, and Piotr Bojanowski.
553 Dinov3, 2025. URL <https://arxiv.org/abs/2508.10104>.

554 Sang Truong, Yuheng Tu, Percy Liang, Bo Li, and Sanmi Koyejo. Reliable and efficient amortized model-based
555 evaluation, 2025. URL <https://arxiv.org/abs/2503.13335>.

556 Sang T. Truong and Sanmi Koyejo. *AI Measurement Science: A Science of Knowing Where AI Thrives, Where*
557 *It Breaks, and How to Respond*. Stanford University, 2026. URL [https://aimslab.stanford.edu/](https://aimslab.stanford.edu/textbook/)
558 [textbook/](https://aimslab.stanford.edu/textbook/).

559 Angelina Wang, Michelle Phan, Daniel E. Ho, and Sanmi Koyejo. Fairness through difference awareness:
560 Measuring desired group discrimination in llms, 2025. URL <https://arxiv.org/abs/2502.01926>.

561 Sheng Zhang, Yanbo Xu, Naoto Usuyama, Hanwen Xu, Jaspreet Bagga, Robert Tinn, Sam Preston, Rajesh Rao,
562 Mu Wei, Naveen Valluri, Cliff Wong, Andrea Tupini, Yu Wang, Matt Mazzola, Swadheen Shukla, Lars Liden,
563 Jianfeng Gao, Angela Crabtree, Brian Piening, Carlo Bifulco, Matthew P. Lungren, Tristan Naumann, Sheng
564 Wang, and Hoifung Poon. Biomedclip: a multimodal biomedical foundation model pretrained from fifteen
565 million scientific image-text pairs, 2025. URL <https://arxiv.org/abs/2303.00915>.